

HACK: Supplementary Material

Additional accuracy results for “HACK: Homomorphic Acceleration via Compression of the Key-Value Cache for Disaggregated LLM Inference” (ACM SIGOPS ATC 2026)

Zeyu Zhang
University of Virginia
Charlottesville, Virginia, USA

Haiying Shen
University of Virginia
Charlottesville, Virginia, USA

Shay Vargaftik
VMware Research by Broadcom
Palo Alto, California, USA

Ran Ben Basat
University College London
London, UK
Broadcom
Palo Alto, California, USA

Michael Mitzenmacher
Harvard University
Boston, Massachusetts, USA

Minlan Yu
Harvard University
Boston, Massachusetts, USA

This document supplements the paper “HACK: Homomorphic Acceleration via Compression of the Key-Value Cache for Disaggregated LLM Inference” (ACM SIGOPS ATC 2026, <https://doi.org/10.1145/3828523.3857108>) with the per-dataset accuracy tables summarized in its Section 6.3. The code of HACK is at <https://github.com/Zeyu-ZEYU/HACK> (archived at <https://doi.org/10.5281/zenodo.22982679>). Table numbers below are local to this document.

1 Additional Accuracy Results

Tables 1–6 report per-dataset accuracy across all evaluated models; table 7 reports additional metrics (perplexity, ROUGE-L, and pass@1/pass@10) for Qwen3 32B.

Table 1. Accuracy (IMDb).

	Mis	Phi	Qwen	Yi	Lla	Fal
Baseline	84.81%	87.84%	93.15%	93.87%	95.73%	85.63%
HACK ($\Pi=32$)	83.79%	87.08%	93.02%	92.82%	94.73%	84.96%
HACK ($\Pi=64$)	83.46%	86.64%	92.48%	92.31%	94.17%	84.72%
CacheGen	83.04%	86.15%	91.19%	91.79%	93.85%	83.94%
KVQuant	82.97%	86.05%	91.10%	91.54%	93.89%	84.12%
HACK ($\Pi=128$)	82.71%	86.07%	90.75%	91.35%	93.76%	84.03%

Table 2. Accuracy (arXiv).

	Mis	Phi	Qwen	Yi	Lla	Fal
Baseline	79.40%	86.35%	88.50%	87.75%	83.79%	79.42%
HACK ($\Pi=32$)	78.53%	85.51%	87.73%	86.69%	82.63%	78.75%
HACK ($\Pi=64$)	78.26%	85.08%	87.25%	86.39%	82.31%	78.28%
CacheGen	77.95%	84.02%	85.74%	85.14%	81.85%	77.74%
KVQuant	77.64%	84.11%	85.50%	85.08%	81.76%	77.72%
HACK ($\Pi=128$)	77.72%	84.49%	86.92%	86.15%	81.97%	77.94%

Table 3. Accuracy (Cocktail).

	Mis	Phi	Qwen	Yi	Lla
Baseline	75.18%	83.92%	89.25%	85.25%	86.39%
HACK ($\Pi=32$)	74.63%	83.05%	88.47%	84.39%	85.54%
HACK ($\Pi=64$)	74.42%	82.66%	88.06%	84.04%	84.87%
CacheGen	73.74%	82.24%	87.61%	83.51%	84.71%
KVQuant	73.72%	82.26%	87.58%	83.42%	84.68%
HACK ($\Pi=128$)	73.81%	82.19%	87.44%	82.86%	84.65%

Table 4. Accuracy (HumanEval).

	Mis	Phi	Qwen	Yi	Lla	Fal
Baseline	89.37%	91.62%	92.17%	90.79%	92.45%	85.21%
HACK ($\Pi=32$)	88.68%	90.50%	91.23%	90.02%	91.28%	84.26%
HACK ($\Pi=64$)	88.19%	90.16%	90.85%	89.63%	91.04%	83.82%
CacheGen	87.73%	89.75%	90.36%	88.94%	90.49%	83.47%
KVQuant	87.64%	89.65%	90.27%	88.93%	90.43%	83.29%
HACK ($\Pi=128$)	87.60%	89.54%	90.08%	88.91%	89.77%	83.19%

Table 5. Accuracy (GSM8K).

	Mis	Phi	Qwen	Yi	Lla	Fal
Baseline	96.1%	94.3%	95.9%	96.2%	95.6%	83.7%
HACK ($\Pi=32$)	95.4%	93.6%	95.0%	95.2%	94.5%	82.8%
HACK ($\Pi=64$)	95.2%	93.2%	94.6%	94.7%	94.1%	82.3%
CacheGen	94.6%	92.8%	94.2%	94.4%	93.7%	81.7%
KVQuant	94.5%	92.7%	94.1%	94.4%	93.5%	81.6%
HACK ($\Pi=128$)	94.2%	92.3%	93.6%	93.8%	93.1%	81.2%

Table 6. Accuracy (RULER-128K).

	Mis	Phi	Qwen	Yi	Lla
Baseline	48.1%	55.7%	85.6%	77.3%	66.6%
HACK ($\Pi=32$)	47.7%	55.2%	85.1%	76.9%	66.2%
HACK ($\Pi=64$)	46.9%	54.7%	84.5%	76.0%	65.4%
CacheGen	46.0%	53.9%	83.7%	75.4%	64.8%
KVQuant	46.2%	53.6%	83.2%	75.5%	64.3%
HACK ($\Pi=128$)	45.1%	52.8%	82.4%	74.2%	63.6%

Additional metrics. Beyond exact match, ROUGE-1, and edit similarity, Table 7 reports perplexity (WikiText-2 [1]), ROUGE-L (arXiv summarization), and pass@1/pass@10 (HumanEval) for Qwen3 32B. HACK ($\Pi=32$) and HACK ($\Pi=64$) track the BF16 baseline most closely and beat CacheGen and KVQuant on all four metrics, with the smaller partition slightly ahead, consistent with the exact-match results above.

Table 7. Additional accuracy metrics for Qwen3 32B.

	WikiText-2	arXiv	HumanEval	
Method	PPL ↓	ROUGE-L	pass@1	pass@10
Baseline	7.6	81.0%	87.0%	93.0%
CacheGen	8.2	78.0%	83.5%	90.5%
KVQuant	8.3	77.8%	83.0%	90.2%
HACK ($\Pi=32$)	7.8	80.2%	86.0%	92.4%
HACK ($\Pi=64$)	7.9	79.6%	85.3%	91.9%

References

- [1] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. 2017. Pointer Sentinel Mixture Models. In *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1609.07843>